

MPS

Lesson 05

Linux

Character Drivers

Today's Lesson

12:1
5



Last Lesson



Presentation of last exercise



Char driver intro and Major/Minor Numbers

12:
50



Break (10 min)



Nodes



File Operations



Device Registration

13:
25



Break (10 min)



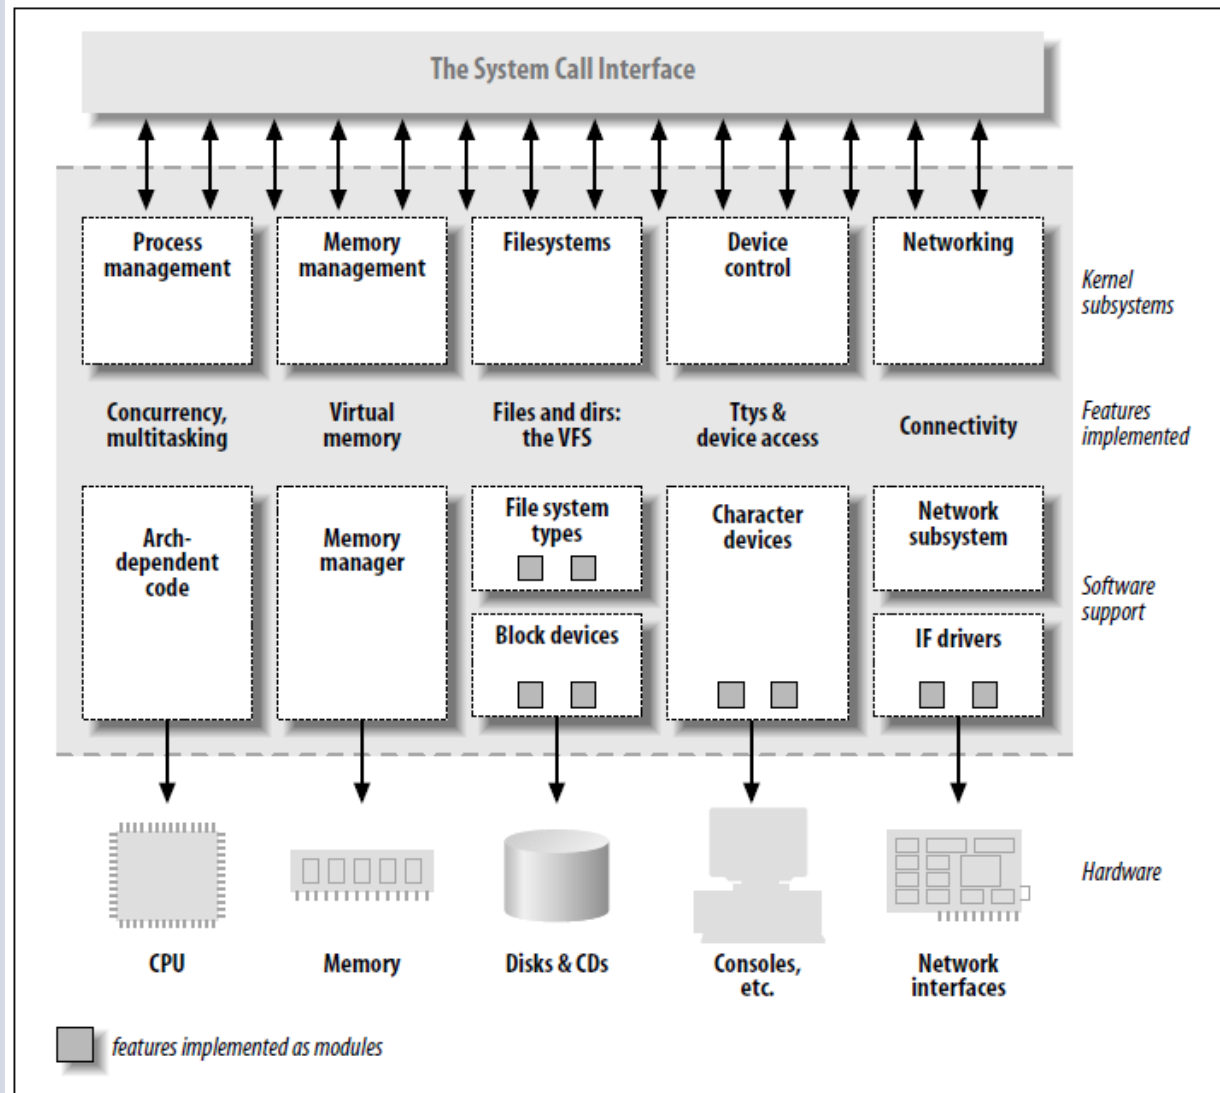
Driver implementation of Fops

13:
40



Exercise: GPIO Linux Device Driver

The Linux Kernel



- A device where
 - Data is accessed as a stream
 - Random read/write to different areas is not required
- Examples
 - Serial ports (RS-232/RS-485), GPIO, A/D Converter, etc
 - Devices not being memory devices or network interfaces



ASR33 Teletype
Origin of the
abbreviation tty

/dev/

```
root@DevKit8000:~# ls -l /dev
crw-rw----    1 root    root    252,    0 Jan  1  1970 DW9710
crw-----    1 root    root      5,    1 Jan 27 10:20 console
crw-rw----    1 root    video   29,    0 Jan  1  1970 fb0
```

Type “tty” to get your current terminal

- Says nothing about the functionality
- Can be allocated statically in the driver or dynamically during module insertion
- Character / Block Devices has each their set of numbers
- Typically one major number per ***driver***

- Specifies the device that uses the driver
- *One minor number per **device***
- Examples
 - Terminal driver
 - Major Number: 4
 - Minor Numbers: 0-xx for tty0, tty1...ttyS2 etc
 - The minor numbers may cover different chips
 - A/D Converter 8-channel (one chip!) driver
 - Major Number: 63
 - Minor Numbers: 0-7
 - Gives us to one (logical) device per channel
 - Adding a chip, could be implemented by adding a major number or just by adding minor numbers.
- Gives us a unified interface to hardware

Major / Minor Numbers

ads7870driver.c

MAJOR = 63
MINOR = 0



```
root@DevKit8000:~# ls -l /dev
crw-rw----  1 root  root   63, 0 Jan  1  1970 ADS7870
```


What Major / Minor Numbers are assigned to the Serial Ports on the Add-on board (ttyS0-ttyS2)?

Static Allocation of Major Numbers

- Drivers for non-general platforms, can use statically allocated numbers
 - This can be the case for a small embedded platforms where you have control of all drivers
- Example:

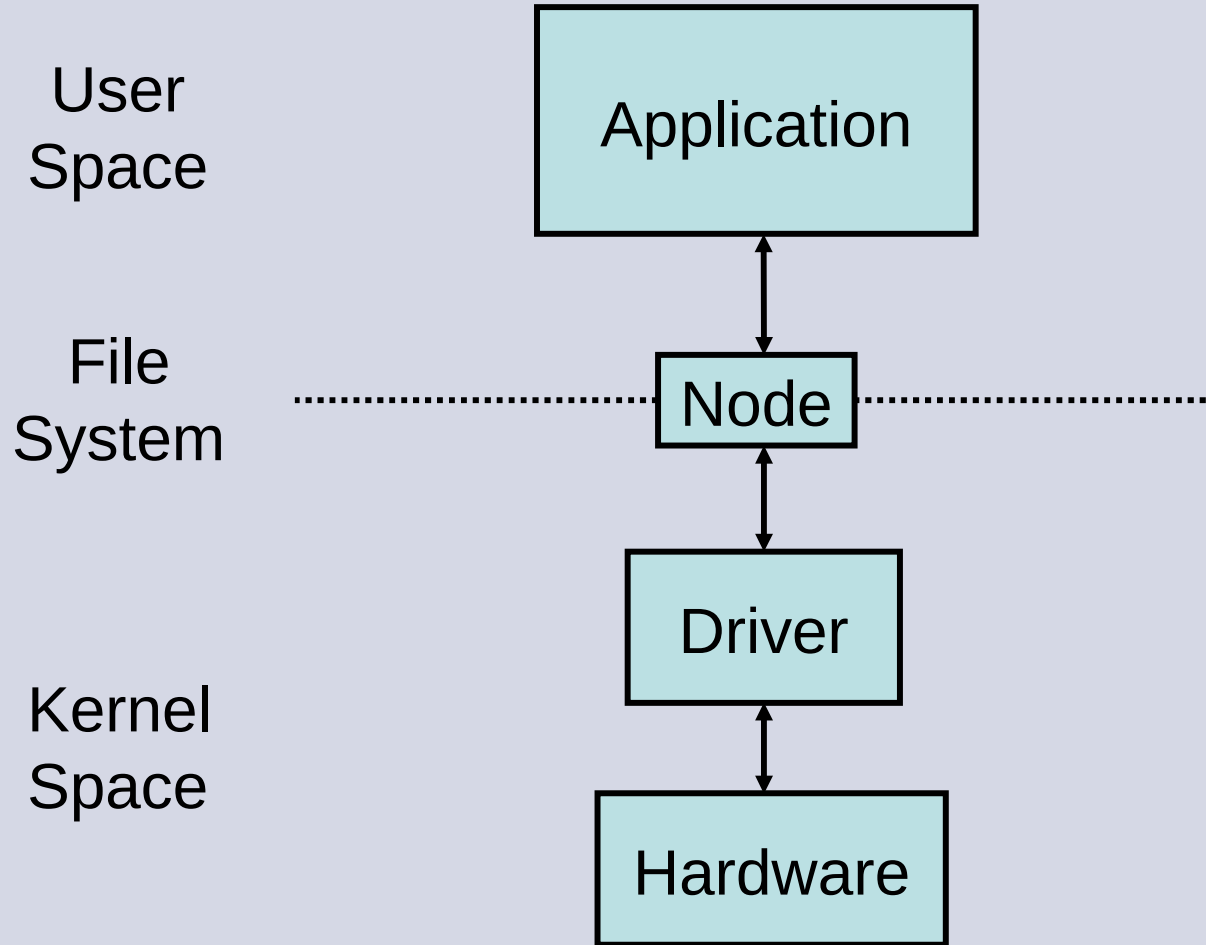
```
#define ADS7870_MAJOR 64
#define ADS7870_MINOR 0
#define ADS7870_CH      8    // 8-ch ADC
...
// Allocating Device Numbers
devno = MKDEV(ADS_7870_MAJOR, ADS7870_MINOR);
err    = register_chrdev_region(devno, ADS7870_CH,
                                "ads7870");

// Freeing Device Numbers
unregister_chrdev_region(devno, ADS7870_CH);
```

- Drivers for general platforms, ex a PC **must** use dynamically allocated numbers
- Example:

```
#define ADS7870_MINOR  0
#define ADS7870_CH      8    // 8-ch ADC
...
// Allocating Device Numbers
err = alloc_chrdev_region(pDev, ADS7870_MINOR,
                          ADS7870_CH, ads7870");

// Freeing Device Numbers
unregister_chrdev_region(devno, ADS7870_CH);
```



- To let the driver appear in the file system, you have to create nodes.

- Examples

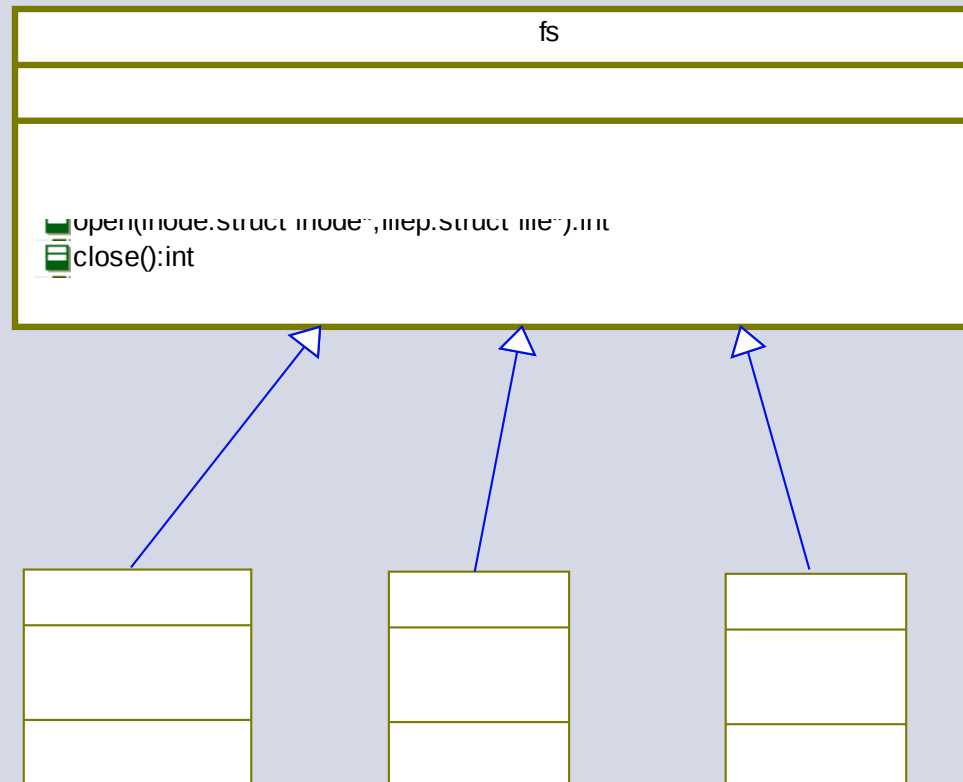
- Driver with statically bound numbers:

```
root@DevKit8000:~# mknod /dev/adc0 c 63 0
```

Desired device name **Char Driver** **Major / Minor**

- Driver with dynamically assigned numbers:
 - Insert the module with “insmod”
 - look up the major number assigned using:
“cat /proc/devices”
 - Create nodes with “mknod” using the major number found
 - See script p.47 in *ldd3*

Linux File Operations (1)



- We must implement the concrete read/write/open/release operations
- Object Oriented C

- We must specify our actual implementation of the file system operations:


```
struct file_operations myDriver_fops = {  
    .owner = THIS_MODULE,           ← attribute  
    .read  = myDriverRead,          ← operations  
    .write = myDriverWrite,  
    .open  = myDriverOpen,  
    .release = myDriverRelease,  
};  
...  
err = myDriverOpen(struct inode *inode,  
                    struct file *filep) {  
    ...  
}
```

Driver Life-Cycle

Module Insertion

File
Operations

Module Extraction



- After allocating major / minor numbers, the driver must be registered in the kernel
- Several helper methods are provided in cdev.h
- For a stand-alone c-dev structure use:

```
#include <linux/cdev.h>

Struct cdev *my_cdev = cdev_alloc();
my_cdev->ops = &my_fops;
```

Device Registration (2)

- For the structure to be located in the driver (typical):

```
#include <linux/cdev.h>

static struct cdev my_cdev;
...
devno = MKDEV(...)
register_chrdev_region(...)

cdev_init(&my_cdev, &my_fops); // Get c-dev
my_cdev.owner = THIS_MODULE;   // Set c-dev owner
my_cdev.ops = &my_fops;        // Set c-dev file op

err = cdev_add(my_cdev, devno, no_devices);
```

- One of the methods inherited from fs is open
- You must implement the concrete tasks:
 - Check for device not ready errors
 - Initialize device if opened for the first time
 - Allocate memory to be used
 - Check if module is already in use (concurrency)

```
int my_cdev_open(struct inode *inode, struct *filep){  
    if (MAJOR(inode->i_rdev) != my_drv_MAJOR)  
        return -ENODEV;  
    if (MINOR(inode->i_rdev) != my_drv_MINOR)  
        return -ENODEV;  
    if (!try_module_get(my_drv_fops.owner))  
        return -ENODEV;  
    return 0; }
```

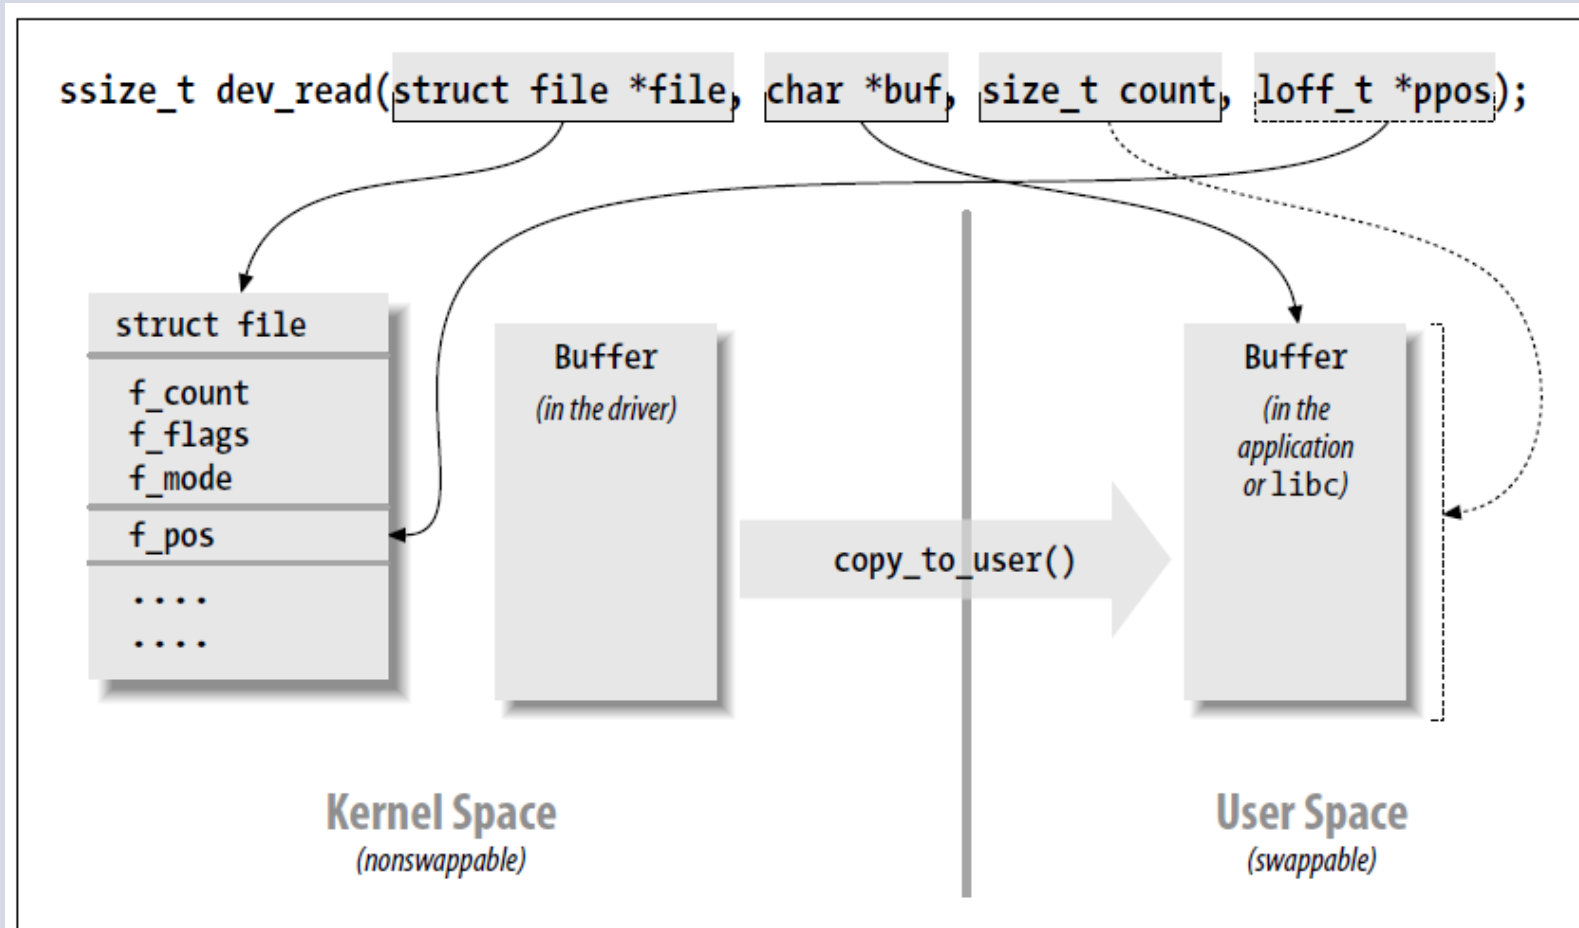
- The Release method is in charge of releasing the module, to make it available for other clients (applications)
- Things to be done during release are:
 - De-allocation of memory allocated during open
 - Shut down the device on the last close
 - Freeing the module semaphore (concurrency)

```
int my_cdev_release(struct inode *inode,
                    struct *filep){
    <... check major/minor ...>
    module_put(my_drv_fops.owner);
    return 0; }
```

User- / Kernel Space Data (1)

- The kernel- and user space has separate memory locations
- User space has only user access, no access to kernel memory. Trying → Page Fault
- User space memory is typically virtual and maybe swappable
- Copy from Kernel- to User Space:
 - unsigned long **copy_to_user**(void __user *to, const void *from, unsigned long count);
- Copy from User- to Kernel Space:
 - unsigned long **copy_from_user**(void *to, const void __user *from, unsigned long count);

User- / Kernel Space Data (2)



- Note that `buf` is a char pointer!

The Read Method

- The “read” method is the driver method called, when an application performs a read (ex: cat /dev/my_dev) on the device
- “read” return the number of bytes read, zero if EOF reached or a negative value on errors

```
ssize_t my_cdev_read(struct file *filep,  
                    char __user *buf,  
                    size_t count, requested  
                    loff_t *f_pos) size  
  
    char my_buf[12] = "hello_world";  
    char my_buf_len = sizeof(my_buf);  
    copy_to_user(buf, my_buf, my_buf_len);  
    *f_pos += my_buf_len;  
    return my_buf_len; }
```

The Write Method

- The “write” method is the driver method called, when an application performs a write (ex: echo “hello” > /dev/my_dev) to the device
- “write” return the number of bytes written, zero if none written or a negative value on errors

```
ssize_t my_cdev_write(struct file *filep,  
                      char __user *buf,  
                      size_t count,  
                      loff_t *f_pos){  
    // my_buf is a local variable  
    copy_from_user(my_buf, buf, count);  
    *f_pos += count;  
    return count; }
```